

Proposal: AI Accelerated Software Development (10-Week Pilot)

Prepared for: DevDatta Halbe (SVP, Product Engineering)

Prepared by: Praxis

Date: 2026-01-20

Executive summary

This proposal describes a 10-week pilot to enable AI-accelerated software delivery for 1-3 development teams. The focus is practical: ship more reliably using agentic coding workflows (plan -> diff -> test -> narrate), under enterprise security constraints, with measurable outcomes.

The pilot is structured as three phases:

- **Fundamentals (Weeks 1-2):** mental model + safety gate (recordable courseware)
- **Labs (Weeks 3-6):** hands-on labs producing PRs, tests, docs, and runbooks
- **Applied (Weeks 7-10):** coached implementation on real Ensemble work + executive readout

Pilot goals and measurable outcomes

Goals by Week 10

- A repeatable, safe workflow for AI-assisted development (tooling + rules + review discipline)
- Teams consistently using an agentic coding method for delivery work
- At least one high-signal improvement per team (capstone PRs) backed by metrics

Target measurable outcomes

- Delivery: PR throughput, time-to-merge, cycle time (ticket -> merge)
- Quality: test/coverage deltas on high-risk areas, reduced review thrash
- Enablement: adoption rate (self-report + signals where available), developer satisfaction

Scope and delivery model

Cohort sizing

- 1-3 cohorts, up to 20 participants per cohort
- Recommended first run: ~15 participants (cap ~20)

Weekly cadence (typical)

- 1 module block (~1 hour)

- 2 labs/workshops (~2 hours each)
- 1 review/Q&A; block (~1 hour)
- short submission/retro (~1 hour)

Total structured time: ~6 hours/week, with Applied work overlapping normal delivery.

Week-by-week (high level)

Week	Theme
1	Fundamentals + safe use kickoff; baseline metrics; enablement discovery
2	Agentic coding foundations; first sandbox PR
3	Lab sprint 1: feature delivery + PR narration workflow
4	Lab sprint 2: testing acceleration + coverage delta reporting
5	Lab sprint 3: docs + architecture extraction (ADR/runbook)
6	Workflow integration into existing SDLC (branching, PR rules, review norms)
7	Coached delivery on real work (safe PR selection + execution)
8	Higher-impact tickets + quality gates
9	Hardening + metrics capture + rollout guidance drafting
10	Capstone + executive readout + next 90-day plan

Deliverables

Labs (Weeks 3-6)

- PR narration templates (what/why/risk/how-tested)
- Architecture README + ADR v1 + dependency/boundary notes
- Risk-ranked backlog of safe PR candidates
- Test plan + coverage/test delta report
- Observability + security drill artifacts (logging schema, correlation IDs, redaction checklist)

Applied (Weeks 7-10)

- Capstone PR package: 1-3 meaningful PRs per team
- Team AI SDLC playbook v1 (workflow, prompt rules, permissions, review gates)
- Executive readout: KPI deltas + rollout recommendations + next 90-day plan

Enablement and security assumptions

Tooling: Claude Code CLI (approved for the pilot)

Environment options (choose early)

- Isolated cloud lab environment (recommended)
- Dev containers on developer machines
- Windows + WSL-based setup

Safety rules (non-negotiable)

- No PHI/PII or sensitive code/data in prompts unless explicitly approved and governed
- Least-privilege tool permissions (read-only first where appropriate)
- Human review required for merges and any production-affecting changes
- Prompt injection awareness: untrusted inputs are isolated and labeled

Roles and responsibilities

Praxis

- Deliver instruction, labs, and coaching
- Provide templates, rubrics, and workflow playbooks
- Support enablement decisions (policy + environment patterns)
- Run measurement and produce executive readout

Ensemble

- Identify 1-3 target teams and participant list

- Provide time protection for participants (8-10 hrs/week)
- Provide access to tooling and a safe lab environment
- Provide KPI sources where available (ADO, dashboards) or agree to manual sampling

Commercials

Hourly pilot (recommended for first run)

Ensemble already has an approved contract vehicle for hourly billing. Praxis proposes delivering the first run under that contract to minimize procurement friction.

- Estimated instructor effort (first pilot): ~150-210 hours total
- Pilot cost: hourly_rate x (150-210)
- Billed time reflects actual delivery effort

List pricing (for future rollouts)

Component	Unit	List price
Fundamentals (Phase 1)	per seat	\$500
Labs (Phase 2)	per seat	\$2,500
Applied (Phase 3)	per cohort	\$20,000

Program total formula (C cohorts, N seats per cohort): $C \times (\$3,000 \times N + \$20,000)$.

Decisions needed to start (end of Week 1)

- Cohort selection (1-3 teams) and participant list
- Environment choice (cloud lab vs dev containers vs WSL)
- Tool access path (Claude Code CLI usage + logging policy)
- Data handling rules (what can/can't enter prompts; redaction expectations)
- KPI access approach (ADO + dashboards vs manual sampling)

Next steps

- Confirm pilot start date and cohort count (1-3)
- Lock environment choice + access plan
- Schedule Week 1 baseline + enablement kickoff
- Publish participant calendar blocks for predictable cadence